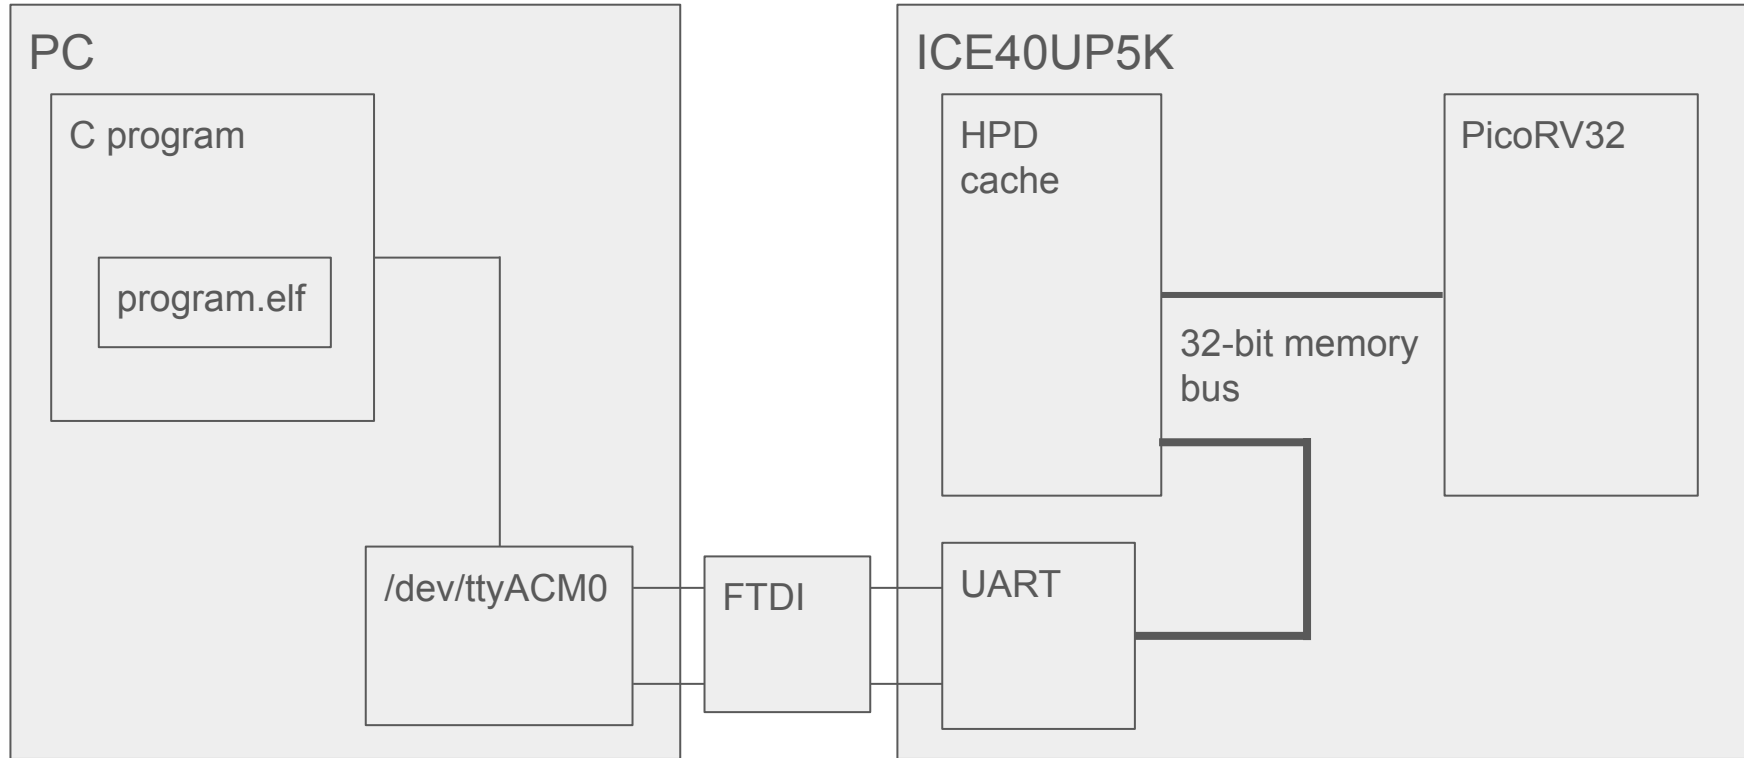


# PicoRV memory over UART

Aditya Bedekar, Lennan Tuffy

# Project description



# UART

Based on the PicoRV's memory requests, send command bytes and either:

- Read: address only
- Write: byte mask, address, and data

To the computer running the uart controller program.

Used a state machine to choose which bytes to send, similar to what HW1 did

Needed extra modifications after adding the cache in order to communicate using the cache-memory interface

# PICORV32

- RISC-V CPU by yosys, works well with low-power FPGAs
- Core: PicoRV32 (4-stage pipeline)
- Target: iCE40UP5K FPGA
- Memory: 128KB (14-bit physical address)

```
// Key parameters from config_pkg.sv
paWidth: 14,           // 128KB memory address space
wordWidth: 32,         // Match PicoRV32's datapath
sets: 64,              // Optimized for iCE40UP5K BRAM
ways: 8,               // Total cache size: 64*8*8*32 = 131,072 bits
```

# PICORV32 - continued

- **Pipeline Optimizations**
- MSHR ( Miss Status Handling Registers) configured for pipeline stages:
- 4 MSHR sets (one per pipeline stage)
- 2 MSHR ways for concurrent misses
- Write buffer: 8 directory entries, 4 data entries
- Memory Interface
- 32-bit address and data width
- Write-back policy for better performance
- Transaction ID width: 4 bits (16 outstanding transactions)

# PICORV32- MSHR

- Purpose: MSHRs track outstanding cache misses to avoid blocking new memory accesses.
- Eg:
  - 1) CPU requests data at address 0X1000, cache miss
  - 2)MSHR store its request, while allowing CPU to continue
  - 3) If another request 0x1004 hits same cache line, MSHR can merge it into the same transaction. In our configuration there are 4 MSHR sets, and 2 MSHR ways which allows for two concurrent cache misses.
- **Write Buffer:** holds write operations before they are sent to main memory.
- 8 directory entries ( to track where data goes), 4 data entities ( to temporarily store data before writing)
- Transaction ID width: used to identify memory requests

# HPD CACHE

- open-source High-Performance, Multi-requester, Out-of-Order L1 Dcache for RISC-V cores and accelerators.
- hpdcache\_lint: 32 bit datapath ,single requester, memory interface signals for read and write operations
- Hpdcache\_wrapper: adapter module between picorv32 and hpdcache

Cache control signals, write buffer management

- Documentation:- <https://github.com/openhwgroup/cv-hpdcache/releases>

# HPD CACHE - continued

- **PMA:** *An underlying system can have multiple physical memory attributes (PMA)*
- **Cacheability:** For a memory space that is segmented ( where segment is defined from base address to end address) some segments are cacheable. The HPD CACHE needs to know which segments are cacheable .
- Eg: for a given read request, should read data be copied into the cache
- Request interface implements an uncacheable bit, when bit is set to 1 , the access is uncacheable.
- Eg: For uncacheable memory, consider I/O device register like a UART serial port.

*To ensure latest value is read from hardware*

		<code>phys_indexed = 1)</code>
<code>core_req_i.pma.uncacheable</code>	Requester	Indicates whether the access needs to be cached (unset) or not (set). Uncacheable accesses are directly forwarded to the memory. It is only valid when using physical indexing ( <code>core_req_i.phys_indexed = 1</code> )
<code>core_req_i.pma.io</code>	Requester	Indicates whether the request targets input/output (IO) periph-



# HPD CACHE - continued

- For the cache memory interface (CMI) , these are a few types of memory request operation types that are supported

## 2.8.1 CMI Type of operation

Table 2.15: Memory request operation types

Mnemonic	Encoding	Type
HPDCACHE_MEM_READ	0b00	Read operation
HPDCACHE_MEM_WRITE	0b01	Write operation
HPDCACHE_MEM_ATOMIC	0b10	Atomic operation

# C program and ELF

Install risc-v gnu toolchain -> picorv recommends the version from 2018

Writing the firmware files in start.S and firmware.c

Compiling and linking to create a .elf file

Reading the .elf file to respond to the core's requests

- ELF headers - elf header, program header, section headers

- Storing the file contents in an accessible way

# Memory map

```
#define OUTPORT 0x10000000

void putc(char ch){
    *((volatile uint32_t*)OUTPORT) = ch;
}
```

The C program interprets writes to 0x10000000 as print statements.

The “virtual” memory doesn’t have space for this address

- Automatically print to stdout if memory is out of bounds

```
char digit_to_char(uint8_t digit){
    char out;
    switch (digit) {
        case 0x0: out = '0'; break;
        case 0x1: out = '1'; break;
        case 0x2: out = '2'; break;
        case 0x3: out = '3'; break;

        case 0x4: out = '4'; break;
        case 0x5: out = '5'; break;
        case 0x6: out = '6'; break;
        case 0x7: out = '7'; break;

        case 0x8: out = '8'; break;
        case 0x9: out = '9'; break;
        case 0xa: out = 'a'; break;
        case 0xb: out = 'b'; break;

        case 0xc: out = 'c'; break;
        case 0xd: out = 'd'; break;
        case 0xe: out = 'e'; break;
        case 0xf: out = 'f'; break;
        default: out = 0;
    }
    return out;
}

void print_hex(unsigned int num, int digits){
    putc('0');
    putc('x');
    for (int i = (4*digits)-4; i >= 0; i -= 4){
        unsigned int mask = 0xf << i;
        char digit = digit_to_char((num & mask) >> i);
        putc(digit);
    }
}
```

# What did we learn?

ELF: writing firmware with riscv-gnu-toolchain, using newlib

Verilator: DPI-C with elf.h - both the c program and tb uses the same file

SystemVerilog language features: structs, sv header files, macros for typedefs

Integration an open-source cache with another open-source risc-v cpu core